

# Q4 - Artificial Neural Networks

---

*Train, validate, and test a neural network model using the dataset in hw3q4.zip, which contains training data (75%), validation data (12.5%), and test data (12.5%). There are two output classes in this data set. You can either choose matlab or a python neural networks package, Keras for this problem.*

## Import Libraries

---

```
1 import matplotlib.pyplot as plt
2 import numpy as np
3 import zipfile
4 import math
5 import sys
6 import os
7 np.random.seed(7)
8 from tensorflow.keras.layers import Activation
9 from tensorflow.keras.layers import Input
10 from tensorflow.keras.layers import Dense
11 from tensorflow.keras import Model
12 import tensorflow as tf
13 tf.set_random_seed(7)
14 %matplotlib inline
```

**(a) Please briefly describe how to construct your working environments (e.g. language, package version, backend for neural networks, installation, etc.) in your report, and write how to execute your codes on 'readme' file.**

---

## Language

- python 3.7

## Prerequisites

- tensorflow 1.8+ (CPU only)

- matplotlib
- numpy

## Installation

Install the above packages using the following commands -

- tensorflow 1.8+ (CPU only) - `pip install tensorflow`
- matplotlib - `pip install matplotlib`
- numpy - `pip install numpy`

## Data

hw3q4.zip should be in the same directory level as this notebook.

## (b) Keras Model

---

### Helper Functions

```
1 def get_data(zipname, filename):
2     zf = zipfile.ZipFile(zipname)
3     data = None
4     with zf.open(filename, 'r') as fp:
5         data = np.asarray([[float(v.strip()) for v in
6             x.decode().strip().split(',') for x in fp.readlines()]])
7     return data
```

```
1 def ANN(n_neurons, output_dims):
2     input_layer = Input(batch_shape = (None, 61), name='input_layer')
3     layer = Dense(n_neurons, name='hidden_layer')(input_layer)
4     layer = Activation(activation='relu', name='relu')(layer)
5     layer = Dense(output_dims, name='output_layer')(layer)
6     output_layer = Activation(activation='sigmoid', name='sigmoid')(layer)
7     model = Model(inputs=input_layer, outputs=output_layer, name='ann')
8     model.compile('adam', 'mse', ['accuracy'])
9     return model
```

### Read Data

```

1 train_x = get_data('Data' + os.sep + 'hw3q4.zip', 'hw3q4' + os.sep +
  'X_train.csv')
2 train_y = get_data('Data' + os.sep + 'hw3q4.zip', 'hw3q4' + os.sep +
  'Y_train.csv')
3 val_x = get_data('Data' + os.sep + 'hw3q4.zip', 'hw3q4' + os.sep +
  'X_val.csv')
4 val_y = get_data('Data' + os.sep + 'hw3q4.zip', 'hw3q4' + os.sep +
  'Y_val.csv')
5 test_x = get_data('Data' + os.sep + 'hw3q4.zip', 'hw3q4' + os.sep +
  'X_test.csv')
6 test_y = get_data('Data' + os.sep + 'hw3q4.zip', 'hw3q4' + os.sep +
  'Y_test.csv')

```

**(1) Construct neural networks using the given training dataset (X train, Y train) using different number of hidden neurons. Set the parameters as follows: activation function for hidden layer='relu', activation for output layer='sigmoid', loss function='mse', metrics='accuracy', epochs=10, batch size=50. For each model, change the number of hidden neurons in the order of 2, 4, 6, 8, 10.**

```

1 hidden_neurons = [2,4,6,8,10]
2 models = [ANN(hidden_neurons[i], train_y.shape[1]) for i in
  range(len(hidden_neurons))]
3 histories = []
4 for i in range(len(models)):
5     print('Training Model with {} neurons in hidden
  layer'.format(hidden_neurons[i]))
6     history = models[i].fit(x=train_x,
7                             y=train_y,
8                             batch_size=50,
9                             epochs=10,
10                            shuffle=True)
11     histories.append(history)
12     print()

```

```

1 Training Model with 2 neurons in hidden layer
2 Epoch 1/10
3 1500/1500 [=====] - 0s 273us/step - loss: 0.2538
  - acc: 0.5333
4 Epoch 2/10

```

```
5 1500/1500 [=====] - 0s 25us/step - loss: 0.2481
  - acc: 0.5533
6 Epoch 3/10
7 1500/1500 [=====] - 0s 24us/step - loss: 0.2432
  - acc: 0.5727
8 Epoch 4/10
9 1500/1500 [=====] - 0s 25us/step - loss: 0.2387
  - acc: 0.5907
10 Epoch 5/10
11 1500/1500 [=====] - 0s 24us/step - loss: 0.2338
  - acc: 0.6087
12 Epoch 6/10
13 1500/1500 [=====] - 0s 25us/step - loss: 0.2285
  - acc: 0.6193
14 Epoch 7/10
15 1500/1500 [=====] - 0s 22us/step - loss: 0.2229
  - acc: 0.6307
16 Epoch 8/10
17 1500/1500 [=====] - 0s 22us/step - loss: 0.2164
  - acc: 0.6347
18 Epoch 9/10
19 1500/1500 [=====] - 0s 22us/step - loss: 0.2095
  - acc: 0.6480
20 Epoch 10/10
21 1500/1500 [=====] - 0s 25us/step - loss: 0.2022
  - acc: 0.6733
22
23 Training Model with 4 neurons in hidden layer
24 Epoch 1/10
25 1500/1500 [=====] - 0s 195us/step - loss: 0.3337
  - acc: 0.4927
26 Epoch 2/10
27 1500/1500 [=====] - 0s 25us/step - loss: 0.3077
  - acc: 0.5280
28 Epoch 3/10
29 1500/1500 [=====] - 0s 24us/step - loss: 0.2827
  - acc: 0.5680
30 Epoch 4/10
31 1500/1500 [=====] - 0s 25us/step - loss: 0.2596
  - acc: 0.5967
32 Epoch 5/10
33 1500/1500 [=====] - 0s 37us/step - loss: 0.2387
  - acc: 0.6307
34 Epoch 6/10
35 1500/1500 [=====] - 0s 27us/step - loss: 0.2213
  - acc: 0.6547
```

```
36 Epoch 7/10
37 1500/1500 [=====] - 0s 24us/step - loss: 0.2066
   - acc: 0.6827
38 Epoch 8/10
39 1500/1500 [=====] - 0s 23us/step - loss: 0.1944
   - acc: 0.7133
40 Epoch 9/10
41 1500/1500 [=====] - 0s 26us/step - loss: 0.1840
   - acc: 0.7360
42 Epoch 10/10
43 1500/1500 [=====] - 0s 24us/step - loss: 0.1753
   - acc: 0.7580
44
45 Training Model with 6 neurons in hidden layer
46 Epoch 1/10
47 1500/1500 [=====] - 0s 209us/step - loss: 0.2943
   - acc: 0.4933
48 Epoch 2/10
49 1500/1500 [=====] - 0s 25us/step - loss: 0.2667
   - acc: 0.5333
50 Epoch 3/10
51 1500/1500 [=====] - 0s 26us/step - loss: 0.2441
   - acc: 0.5673
52 Epoch 4/10
53 1500/1500 [=====] - 0s 25us/step - loss: 0.2260
   - acc: 0.6000
54 Epoch 5/10
55 1500/1500 [=====] - 0s 25us/step - loss: 0.2110
   - acc: 0.6453
56 Epoch 6/10
57 1500/1500 [=====] - 0s 27us/step - loss: 0.1988
   - acc: 0.6727
58 Epoch 7/10
59 1500/1500 [=====] - 0s 26us/step - loss: 0.1883
   - acc: 0.7087
60 Epoch 8/10
61 1500/1500 [=====] - 0s 27us/step - loss: 0.1791
   - acc: 0.7387
62 Epoch 9/10
63 1500/1500 [=====] - 0s 31us/step - loss: 0.1711
   - acc: 0.7573
64 Epoch 10/10
65 1500/1500 [=====] - 0s 23us/step - loss: 0.1643
   - acc: 0.7760
66
67 Training Model with 8 neurons in hidden layer
```

```
68 Epoch 1/10
69 1500/1500 [=====] - 0s 214us/step - loss: 0.2895
   - acc: 0.5273
70 Epoch 2/10
71 1500/1500 [=====] - 0s 25us/step - loss: 0.2614
   - acc: 0.5653
72 Epoch 3/10
73 1500/1500 [=====] - 0s 26us/step - loss: 0.2389
   - acc: 0.6087
74 Epoch 4/10
75 1500/1500 [=====] - 0s 27us/step - loss: 0.2202
   - acc: 0.6460
76 Epoch 5/10
77 1500/1500 [=====] - 0s 27us/step - loss: 0.2051
   - acc: 0.6740
78 Epoch 6/10
79 1500/1500 [=====] - 0s 29us/step - loss: 0.1924
   - acc: 0.7073
80 Epoch 7/10
81 1500/1500 [=====] - 0s 28us/step - loss: 0.1815
   - acc: 0.7347
82 Epoch 8/10
83 1500/1500 [=====] - 0s 25us/step - loss: 0.1725
   - acc: 0.7533
84 Epoch 9/10
85 1500/1500 [=====] - 0s 25us/step - loss: 0.1647
   - acc: 0.7680
86 Epoch 10/10
87 1500/1500 [=====] - 0s 23us/step - loss: 0.1580
   - acc: 0.7747
88
89 Training Model with 10 neurons in hidden layer
90 Epoch 1/10
91 1500/1500 [=====] - 0s 268us/step - loss: 0.2728
   - acc: 0.5547
92 Epoch 2/10
93 1500/1500 [=====] - 0s 28us/step - loss: 0.2387
   - acc: 0.6080
94 Epoch 3/10
95 1500/1500 [=====] - 0s 26us/step - loss: 0.2110
   - acc: 0.6700
96 Epoch 4/10
97 1500/1500 [=====] - 0s 25us/step - loss: 0.1890
   - acc: 0.7140
98 Epoch 5/10
```

```

99 1500/1500 [=====] - 0s 29us/step - loss: 0.1722
    - acc: 0.7473
100 Epoch 6/10
101 1500/1500 [=====] - 0s 22us/step - loss: 0.1596
    - acc: 0.7687
102 Epoch 7/10
103 1500/1500 [=====] - 0s 29us/step - loss: 0.1499
    - acc: 0.7880
104 Epoch 8/10
105 1500/1500 [=====] - 0s 22us/step - loss: 0.1425
    - acc: 0.8053
106 Epoch 9/10
107 1500/1500 [=====] - 0s 28us/step - loss: 0.1364
    - acc: 0.8220
108 Epoch 10/10
109 1500/1500 [=====] - 0s 26us/step - loss: 0.1315
    - acc: 0.8240

```

**(2) Validate each neural network using the given validation dataset (X val, Y val). The validation accuracy is used to determine how many number of hidden neurons are optimal for this problem.**

```

1 eval_results = []
2 for i in range(len(models)):
3     print('Evaluating Model with {} neurons in hidden
    layer'.format(hidden_neurons[i]))
4     eval_result = models[i].evaluate(x=val_x, y=val_y, batch_size=50)
5     print('Mean Squared Error - {:.4f}'.format(eval_result[0]))
6     print('Accuracy - {:.4f}'.format(eval_result[1]))
7     eval_results.append(eval_result)
8     print()

```

```

1 Evaluating Model with 2 neurons in hidden layer
2 250/250 [=====] - 0s 360us/step
3 Mean Squared Error - 0.1976
4 Accuracy - 0.7400
5
6 Evaluating Model with 4 neurons in hidden layer
7 250/250 [=====] - 0s 346us/step
8 Mean Squared Error - 0.1650

```

```

9 Accuracy - 0.7600
10
11 Evaluating Model with 6 neurons in hidden layer
12 250/250 [=====] - 0s 372us/step
13 Mean Squared Error - 0.1585
14 Accuracy - 0.7920
15
16 Evaluating Model with 8 neurons in hidden layer
17 250/250 [=====] - 0s 341us/step
18 Mean Squared Error - 0.1595
19 Accuracy - 0.7440
20
21 Evaluating Model with 10 neurons in hidden layer
22 250/250 [=====] - 0s 318us/step
23 Mean Squared Error - 0.1238
24 Accuracy - 0.8280

```

**(c) Plot a figure, where the horizontal x-axis is the number of hidden neurons, and the vertical y-axis is the accuracy. Please plot both training and validation accuracy in your figure. (Note that the exact accuracy could be slightly different according to your working environments, however you can analyze the trend.**

```

1 fig, ax = plt.subplots(nrows=1, ncols=2, figsize=(16,5))
2 ax[0].set_title('Hidden Number of Neurons vs Training Accuracy')
3 ax[0].set_ylabel('Training Accuracy')
4 ax[0].set_xlabel('Hidden Number of Neurons')
5 ax[0].bar(hidden_neurons,
6           [histories[i].history['acc'][-1] for i in
7            range(len(hidden_neurons))],
8           width=1,
9           color='orange')
10 for i in range(len(hidden_neurons)):
11     ax[0].annotate('{:.4f}'.format(histories[i].history['acc'][-1]),
12                   (hidden_neurons[i]-0.4, histories[i].history['acc']
13                    [-1]-0.05))
14
15 ax[1].set_title('Hidden Number of Neurons vs Validation Accuracy')



```



```

14 ax[1].set_ylabel('Validation Accuracy')
15 ax[1].set_xlabel('Hidden Number of Neurons')
16 ax[1].bar(hidden_neurons,
17           [eval_results[i][1] for i in range(len(hidden_neurons))],
18           width=1,
19           color='gray')
20 for i in range(len(hidden_neurons)):
21     ax[1].annotate('{:.4f}'.format(eval_results[i][1]),
22                   (hidden_neurons[i]-0.4, eval_results[i][1]-0.05))

```



**(d) Provide a simple analysis about your results and choose the optimal number of hidden neuron from the analysis.**

```

1 model = models[4]

```

Based on the above plots, we can clearly see that we get the best validation accuracy when number of neurons in hidden layer is 10. Thus we select the model with 10 neurons in hidden layers.

**(e) Report the test accuracy using the given test dataset (X test, Y test) on the neural network with the optimal number of hidden neurons.**

```

1 print('Testing Model with 10 neurons in hidden layer')
2 eval_result = model.evaluate(x=test_x, y=test_y, batch_size=50)
3 print('Mean Squared Error - {:.4f}'.format(eval_result[0]))
4 print('Accuracy - {:.4f}'.format(eval_result[1]))
5 print()

```

```
1 | Testing Model with 10 neurons in hidden layer
2 | 250/250 [=====] - 0s 19us/step
3 | Mean Squared Error - 0.1594
4 | Accuracy - 0.7880
```

Test Accuracy is 0.7880 with 10 neurons in Hidden Layer.