

# Q6 - SVM Programming

*In this question, you will employ SVM to solve a classification problem for the provided data file "hw3q6.csv". Each row in the data file indicates a sample. The first 12 columns are features and the last column "Class" indicates the label, with 1 and 0 indicating the positive and negative samples, respectively.*

## Import Libraries

```
1 from sklearn.metrics import accuracy_score, f1_score, precision_score,
  recall_score
2 from sklearn.model_selection import train_test_split, GridSearchCV
3 from sklearn.svm import SVC
4
5 import matplotlib.pyplot as plt
6 import numpy as np
7 import sys
8 import os
9
10 %matplotlib inline
```

## (a) Load data. Report the size of positive and negative samples in dataset

```
1 def data_and_headers(filename):
2     data = None
3     with open(filename) as fp:
4         data = [x.strip().split(',') for x in fp.readlines()]
5     headers = data[0]
6     headers = np.asarray(headers)
7     class_field = len(headers) - 1
8     data_x = [[float(x[i]) for i in range(class_field)] for x in data[1:]]
9     data_x = np.asarray(data_x)
10    data_y = [[int(x[i]) for i in range(class_field, class_field + 1)] for
11              x in data[1:]]
12    data_y = np.asarray(data_y)
13    return headers, data_x, data_y
```

```
1 headers, features_x, labels_y = data_and_headers('Data' + os.sep +  
  'hw3q6.csv')
```

```
1 print('Data')
2 print('Number of features - ' + str(features_x.shape[1]))
3 print('Total Number of observations - ' + str(features_x.shape[0]))
4 print('Number of Positive Samples - ' + str(labels_y[labels_y==1].shape[0]))
5 print('Number of Negative Samples - ' + str(labels_y[labels_y==0].shape[0]))
6 print()
```

```
1 Data
2 Number of features - 12
3 Total Number of observations - 200
4 Number of Positive Samples - 90
5 Number of Negative Samples - 110
```

[illegible]

```

1 print('Training Data')
2 print('Number of features - ' + str(train_x.shape[1]))
3 print('Total Number of observations - ' + str(train_x.shape[0]))
4 print('Number of Positive Samples - ' + str(train_y[train_y==1].shape[0]))
5 print('Number of Negative Samples - ' + str(train_y[train_y==0].shape[0]))
6 print()
7 print('Testing Data')
8 print('Number of features - ' + str(test_x.shape[1]))
9 print('Total Number of observations - ' + str(test_x.shape[0]))
10 print('Number of Positive Samples - ' + str(test_y[test_y==1].shape[0]))
11 print('Number of Negative Samples - ' + str(test_y[test_y==0].shape[0]))
12 print()

```

```

1 Training Data
2 Number of features - 12
3 Total Number of observations - 150
4 Number of Positive Samples - 67
5 Number of Negative Samples - 83
6
7 Testing Data
8 Number of features - 12
9 Total Number of observations - 50
10 Number of Positive Samples - 23
11 Number of Negative Samples - 27

```

**(c) Take SVM with linear kernel as classifier (third-party packages are allowed to use) and set the regularization parameter C as: [0.1, 0.5, 1, 5, 10, 50, 100], respectively. For each value of C, train a SVM classifier with the training data and get the number of support vectors (SVs). Generate a plot with C as the horizontal axis and number of SVs as the vertical axis. Give a brief analysis for the plot.**

```

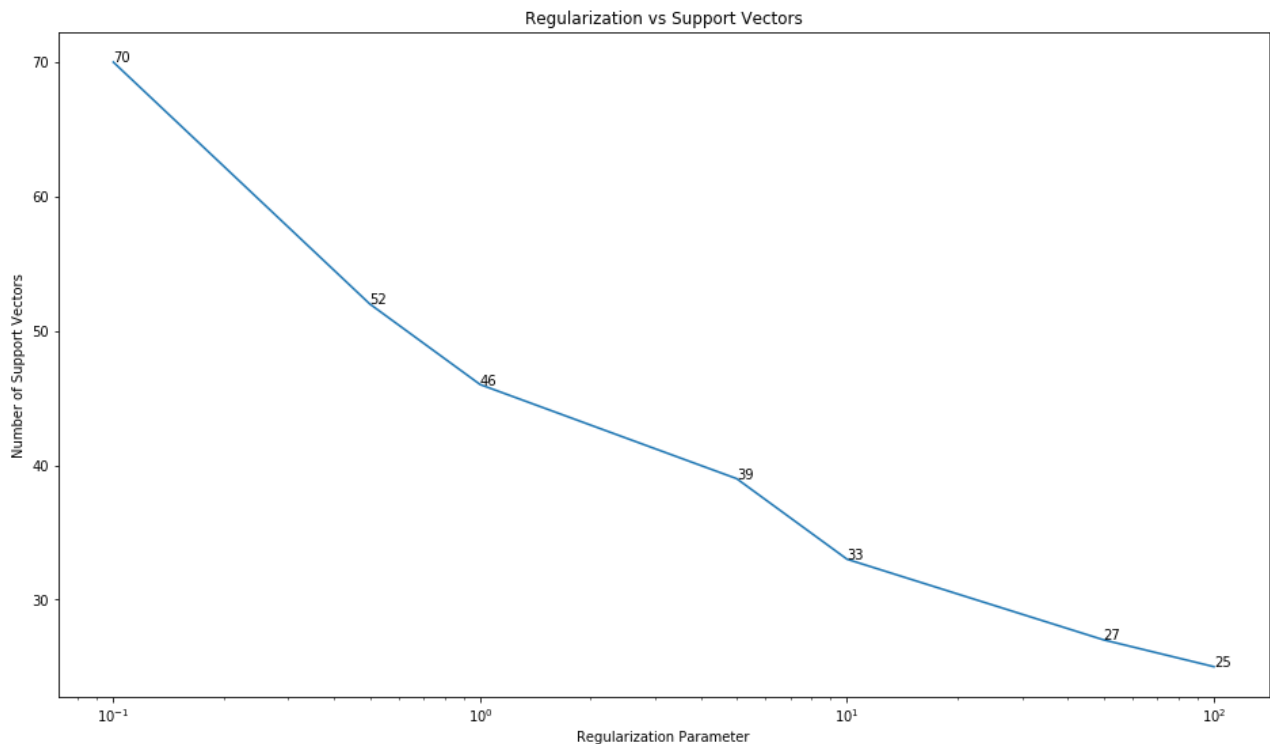
1 C = [0.1, 0.5, 1, 5, 10, 50, 100]
2 models = [SVC(C[i], 'linear', random_state=10).fit(train_x, train_y.flatten())
             for i in range(len(C))]

```

```

1 fig, ax = plt.subplots(figsize=(16,9))
2 ax.set_title('Regularization vs Support Vectors')
3 ax.set_ylabel('Number of Support Vectors')
4 ax.set_xlabel('Regularization Parameter')
5 ax.set_xscale('log')
6 plty = [x.support_.shape[0] for x in models]
7 #ax.bar(C, plty, color='orange')
8 ax.plot(C, plty)
9 for i, txt in enumerate(C):
10     ax.annotate(plty[i], (C[i], plty[i]))
11 plt.show()

```



**(d) Compare 4 different kernel functions, including linear, polynomial, radial basic function (Gaussian kernel), and sigmoid kernel. Make a table to record the accuracy, precision, recall and f-measure of the classification results for the 4 kernel functions. Try to tune the parameters via grid search and report your best results with the optimal parameters. Based on the results, which kernel function will you choose?**

# Linear Kernel

```
1 param_grid = [  
2     {'C':[0.001, 0.005, 0.01, 0.05, 0.1, 0.5, 1, 5, 10, 50, 100],  
3     'kernel':['linear']}]  
4 ]  
5  
6 lsvc = SVC(random_state=10)  
7 linear_clf = GridSearchCV(lsvc, param_grid,  
8     ['f1', 'accuracy', 'recall', 'precision'],  
9     cv=5, refit='accuracy', verbose=1, n_jobs=4)  
10 linear_clf.fit(train_x, train_y.flatten())  
11 y_pred = linear_clf.predict(test_x)
```

```
1 Fitting 5 folds for each of 11 candidates, totalling 55 fits  
2 [Parallel(n_jobs=4)]: Using backend LokyBackend with 4 concurrent workers.  
3 [Parallel(n_jobs=4)]: Done 55 out of 55 | elapsed: 0.3s finished
```

```
1 print('Best Parameters - '  
2 print(linear_clf.best_params_)  
3 print()  
4 print('Linear Kernel Classification Results on Test Set with Optimal  
5 Parameters')  
6 print('Accuracy - \t{:.4f}'.format(accuracy_score(test_y.flatten(),  
7 y_pred)))  
8 print('F1 Score - \t{:.4f}'.format(f1_score(test_y.flatten(), y_pred)))  
9 print('Precision - \t{:.4f}'.format(precision_score(test_y.flatten(),  
10 y_pred)))  
11 print('Recall - \t{:.4f}'.format(recall_score(test_y.flatten(), y_pred)))
```

```
1 Best Parameters -  
2 {'C': 10, 'kernel': 'linear'}  
3  
4 Linear Kernel Classification Results on Test Set with Optimal Parameters  
5 Accuracy - 0.9600  
6 F1 Score - 0.9565  
7 Precision - 0.9565  
8 Recall - 0.9565
```

# Radial Basis Function

```

1 param_grid = [
2     {'C':[0.001, 0.005, 0.01, 0.05, 0.1, 0.5, 1, 5, 10, 50, 100],
3      'kernel':['rbf'],
4      'gamma': [0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,1.0]}
5 ]
6
7 rsvc = SVC(random_state=10)
8 radial_clf = GridSearchCV(rsvc, param_grid,
9                           ['f1', 'accuracy', 'recall', 'precision'],
10                          cv=5, refit='accuracy', verbose=1, n_jobs=4)
11 radial_clf.fit(train_x, train_y.flatten())
12 y_pred = radial_clf.predict(test_x)

```

```

1 Fitting 5 folds for each of 110 candidates, totalling 550 fits
2 [Parallel(n_jobs=4)]: Using backend LokyBackend with 4 concurrent workers.
3 [Parallel(n_jobs=4)]: Done 550 out of 550 | elapsed: 1.4s finished

```

```

1 print('Best Parameters - ')
2 print(radial_clf.best_params_)
3 print()
4 print('Radial Basis Function Kernel Classification Results on Test Set with
5 Optimal Parameters')
6 print('Accuracy - \t{:.4f}'.format(accuracy_score(test_y.flatten(),
7 y_pred)))
8 print('F1 Score - \t{:.4f}'.format(f1_score(test_y.flatten(), y_pred)))
9 print('Precision - \t{:.4f}'.format(precision_score(test_y.flatten(),
10 y_pred)))
11 print('Recall - \t{:.4f}'.format(recall_score(test_y.flatten(), y_pred)))

```

```

1 Best Parameters -
2 {'C': 1, 'gamma': 0.2, 'kernel': 'rbf'}
3
4 Radial Basis Function Kernel Classification Results on Test Set with Optimal
5 Parameters
6 Accuracy - 0.9800
7 F1 Score - 0.9778
8 Precision - 1.0000
9 Recall - 0.9565

```

## Polynomial Kernel

```

1 param_grid = [
2     {'C':[0.001, 0.005, 0.01, 0.05, 0.1, 0.5, 1, 5, 10, 50, 100],
3      'kernel':['poly'],
4      'degree':[2,3,4,5,6],
5      'gamma': [0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,1.0],
6      'coef0': [0.1,0.5,1.0,1.5,2.0,5.0,10.0,20.0]}
7 ]
8
9 psvc = SVC(random_state=10)
10 poly_clf = GridSearchCV(psvc, param_grid,
11                          ['f1', 'accuracy', 'recall', 'precision'],
12                          cv=5, refit='accuracy', verbose=1, n_jobs=4)
13 poly_clf.fit(train_x, train_y.flatten())
14 y_pred = poly_clf.predict(test_x)

```

```

1 Fitting 5 folds for each of 4400 candidates, totalling 22000 fits
2 [Parallel(n_jobs=4)]: Using backend LokyBackend with 4 concurrent workers.
3 [Parallel(n_jobs=4)]: Done 756 tasks      | elapsed:    2.2s
4 [Parallel(n_jobs=4)]: Done 4056 tasks     | elapsed:    9.3s
5 [Parallel(n_jobs=4)]: Done 9556 tasks     | elapsed:   20.5s
6 [Parallel(n_jobs=4)]: Done 17256 tasks    | elapsed:   34.6s
7 [Parallel(n_jobs=4)]: Done 22000 out of 22000 | elapsed:   43.7s finished

```

```

1 print('Best Parameters - ')
2 print(poly_clf.best_params_)
3 print()
4 print('Polynomial Kernel Classification Results on Test Set with Optimal
5 Parameters')
6 print('Accuracy - \t{:.4f}'.format(accuracy_score(test_y.flatten(),
7 y_pred)))
8 print('F1 Score - \t{:.4f}'.format(f1_score(test_y.flatten(), y_pred)))
9 print('Precision - \t{:.4f}'.format(precision_score(test_y.flatten(),
10 y_pred)))
11 print('Recall - \t{:.4f}'.format(recall_score(test_y.flatten(), y_pred)))

```

```
1 Best Parameters -
2 {'C': 0.001, 'coef0': 5.0, 'degree': 4, 'gamma': 0.8, 'kernel': 'poly'}
3
4 Polynomial Kernel Classification Results on Test Set with Optimal Parameters
5 Accuracy - 0.9600
6 F1 Score - 0.9565
7 Precision - 0.9565
8 Recall - 0.9565
```

## Sigmoid Kernel

```
1 param_grid = [
2     {'C':[0.001, 0.005, 0.01, 0.05, 0.1, 0.5, 1, 5, 10, 50, 100],
3      'kernel':['sigmoid'],
4      'gamma': [0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,1.0],
5      'coef0': [0.1,0.5,1.0,1.5,2.0,5.0,10.0,20.0]}
6 ]
7
8 ssvc = SVC(random_state=10)
9 sig_clf = GridSearchCV(ssvc, param_grid,
10                        ['f1', 'accuracy', 'recall', 'precision'],
11                        cv=5, refit='accuracy', verbose=1, n_jobs=4)
12 sig_clf.fit(train_x, train_y.flatten())
13 y_pred = sig_clf.predict(test_x)
```

```
1 Fitting 5 folds for each of 880 candidates, totalling 4400 fits
2 [Parallel(n_jobs=4)]: Using backend LokyBackend with 4 concurrent workers.
3 [Parallel(n_jobs=4)]: Done 416 tasks      | elapsed: 4.6s
4 [Parallel(n_jobs=4)]: Done 2216 tasks    | elapsed: 15.1s
5 [Parallel(n_jobs=4)]: Done 4400 out of 4400 | elapsed: 20.9s finished
```



```

1 print('Best Parameters - ')
2 print(sig_clf.best_params_)
3 print()
4 print('Sigmoid Kernel Classification Results on Test Set with Optimal
Parameters')
5 print('Accuracy - \t{:.4f}'.format(accuracy_score(test_y.flatten(),
y_pred)))
6 print('F1 Score - \t{:.4f}'.format(f1_score(test_y.flatten(), y_pred)))
7 print('Precision - \t{:.4f}'.format(precision_score(test_y.flatten(),
y_pred)))
8 print('Recall - \t{:.4f}'.format(recall_score(test_y.flatten(), y_pred)))

```

```

1 Best Parameters -
2 {'C': 0.05, 'coef0': 0.1, 'gamma': 0.3, 'kernel': 'sigmoid'}
3
4 Sigmoid Kernel Classification Results on Test Set with Optimal Parameters
5 Accuracy - 0.9000
6 F1 Score - 0.9020
7 Precision - 0.8214
8 Recall - 1.0000

```

## Results

Kernel with Tuned Hyperparameter	Accuracy	F1-score	Precision	Recall
<b>Linear</b> {'C': 10}	0.9600	0.9565	0.9565	0.9565
<b>RBF</b> {'C': 1, 'gamma': 0.2}	0.9800	0.9778	1.0000	0.9565
<b>Polynomial</b> {'C': 0.001, 'coef0': 5.0, 'degree': 4, 'gamma': 0.8}	0.9600	0.9565	0.9565	0.9565
<b>Sigmoid</b> {'C': 0.05, 'coef0': 0.1, 'gamma': 0.3}	0.9000	0.9020	0.8214	1.0000

Based on the above results, we will choose the RBF kernel, with C=1 and gamma=0.2 as the kernel function. This is because it is giving the best score across 3 categories - accuracy, f1 and precision.